

Information Management & Systems Engineering

## **Milestone 2: Migration from a relational DBMS to NoSQL**

**Fabian Frauenberger**

11807201

a11807201@unet.univie.ac.at

**Samuel Mitterrutzner**

11911464

a11911464@unet.univie.ac.at

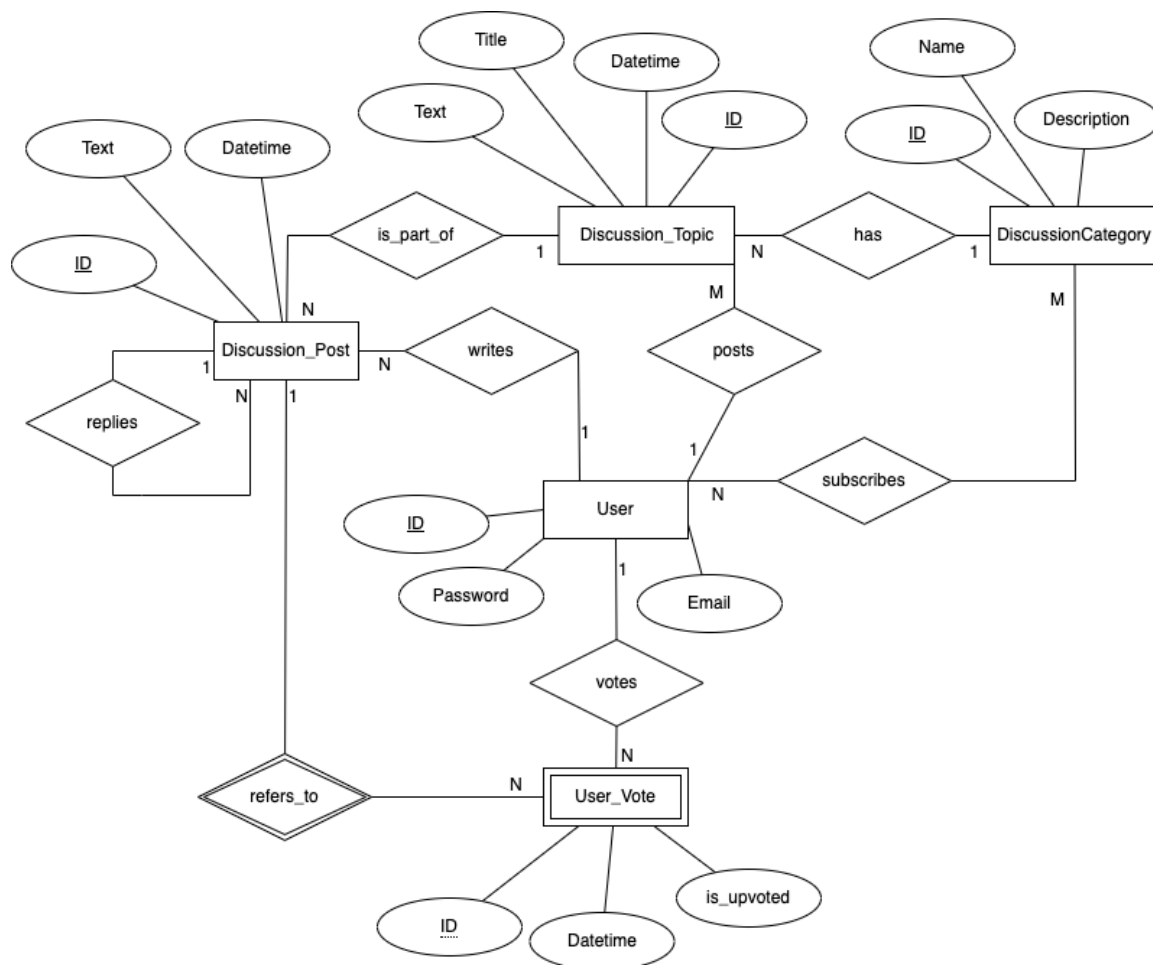
**June 23rd 2022**

## Configuration of Infrastructure

To access the web application, go to the folder submitted and run the docker-compose up command in a CLI. This will launch all the required containers. Then open <http://localhost:8080> or <https://localhost> in your browser. Because we used a self-signed certificate, you may get a certificate warning when using https.

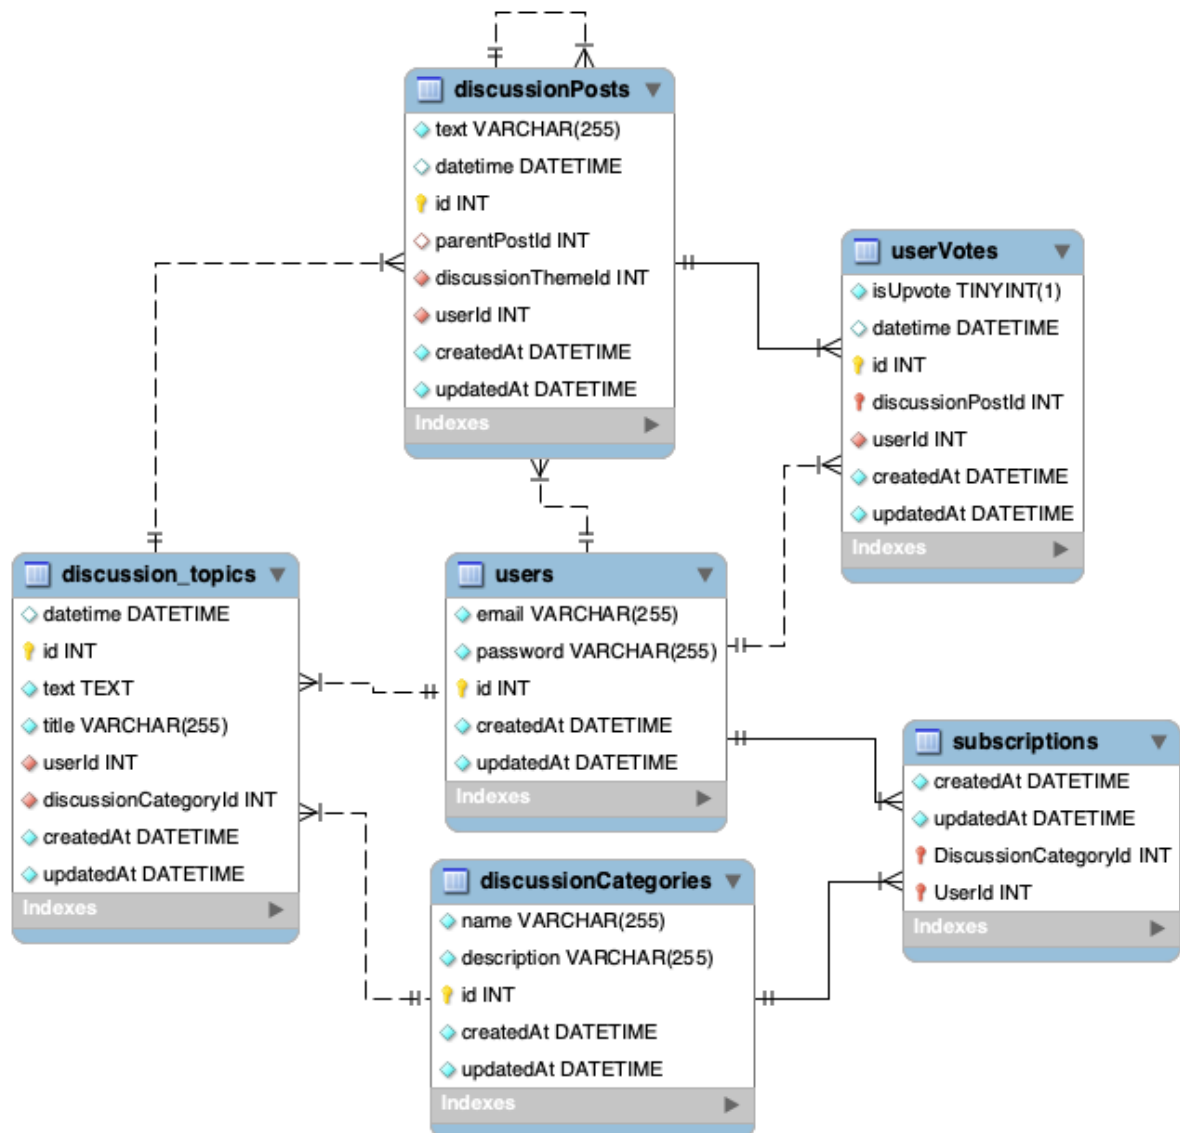
## RDBMS Design

### Entity Relationship Diagram



The idea of our project is to implement an online discussion platform, where users can write anonymously. In comparison to Milestone 1 the username has been removed and only the email (that is not shown publicly) is stored in the user table. Users can create new topics or write posts about topics. A Voting system for posts is another way to agree with other people's opinions.

## Logical Database Design Graph



## Logical Database Design

**users** (id, email, password)

PK: id

**discussionCategories** (id, name, description)

PK: name

**subscriptions** (DiscussionCategoryId, UserId)

PK: DiscussionCategoryId, UserId

FK: subscriptions.DiscussionCategoryId <> discussionCategories.id

subscriptions.UserId <> users.id

**discussion\_topics** (id, title, text, datetime, *discussionCategoryId*, *userId*)

PK: id

FK: discussion\_topics.discussionCategoryId <> discussionCategories.id

discussion\_topics.userId <> users.id

**discussionPosts** (id, text, datetime, *discussionThemeld*, *userId*, *parentPostId*)

PK: id

FK: discussionPosts.discussionThemeld <> discussion\_topics.id

discussionPosts.userId <> users.id

discussionPosts.parentPostId <> discussionPosts.id

**userVotes** (id, discussionPostId, isUpvote, datetime, *userId*)

PK: id, discussionPostId

FK: userVotes.discussionPostId <> discussionPosts.id

userVotes.userId <> users.id

## Physical Database Design

```
CREATE TABLE users (  
    email      VARCHAR(255) NOT NULL UNIQUE,  
    password   VARCHAR(255) NOT NULL,  
    id         INTEGER auto_increment,  
    createdat  DATETIME NOT NULL,  
    updatedat  DATETIME NOT NULL,  
    PRIMARY KEY (id)  
)  
  
CREATE TABLE discussioncategories (  
    name        VARCHAR(255) NOT NULL UNIQUE,  
    description  VARCHAR(255) NOT NULL,  
    id          INTEGER auto_increment,  
    createdat   DATETIME NOT NULL,  
    updatedat   DATETIME NOT NULL,  
    PRIMARY KEY (id)  
)  
  
CREATE TABLE subscriptions (  
    createdat      DATETIME NOT NULL,  
    updatedat      DATETIME NOT NULL,  
    discussioncategoryid INTEGER,  
    userid         INTEGER,  
    PRIMARY KEY (discussioncategoryid, userid),  
    FOREIGN KEY (discussioncategoryid) REFERENCES discussioncategories (id) ON  
    DELETE CASCADE ON UPDATE CASCADE,  
    FOREIGN KEY (userid) REFERENCES users (id) ON DELETE CASCADE ON UPDATE  
    CASCADE  
)
```

```

CREATE TABLE discussion_topics (

    datetime            DATETIME,

    id                  INTEGER auto_increment,

    text                TEXT NOT NULL,

    title               VARCHAR(255) NOT NULL,

    userid              INTEGER NOT NULL,

    discussioncategoryid INTEGER NOT NULL,

    createdat           DATETIME NOT NULL,

    updatedat           DATETIME NOT NULL,

    PRIMARY KEY (id),

    FOREIGN KEY (userid) REFERENCES users (id) ON DELETE CASCADE ON UPDATE

    CASCADE,

    FOREIGN KEY (discussioncategoryid) REFERENCES discussioncategories (id) ON

    DELETE CASCADE ON UPDATE CASCADE

)

```

```

CREATE TABLE discussionposts (

    text                VARCHAR(255) NOT NULL,

    datetime            DATETIME,

    id                  INTEGER auto_increment,

    parentpostid        INTEGER,

    discussionthemeid    INTEGER NOT NULL,

    userid              INTEGER NOT NULL,

    createdat           DATETIME NOT NULL,

    updatedat           DATETIME NOT NULL,

    PRIMARY KEY (id),

    FOREIGN KEY (parentpostid) REFERENCES discussionposts (id) ON DELETE SET

    NULL ON UPDATE CASCADE,

    FOREIGN KEY (discussionthemeid) REFERENCES discussion_topics (id) ON DELETE

    no action ON UPDATE CASCADE,

    FOREIGN KEY (userid) REFERENCES users (id) ON DELETE no action ON UPDATE

```

```
CASCADE  
  
)
```

```
CREATE TABLE uservotes (  
  
    isupvote          TINYINT(1) NOT NULL,  
  
    datetime          DATETIME,  
  
    id                INTEGER auto_increment,  
  
    discussionpostid  INTEGER NOT NULL,  
  
    userid            INTEGER NOT NULL,  
  
    createdat         DATETIME NOT NULL,  
  
    updatedat         DATETIME NOT NULL,  
  
    PRIMARY KEY (id, discussionpostid),  
  
    FOREIGN KEY (discussionpostid) REFERENCES discussionposts (id) ON DELETE no  
action ON UPDATE CASCADE,  
  
    FOREIGN KEY (userid) REFERENCES users (id) ON DELETE no action ON UPDATE  
CASCADE  
  
)
```

# NoSQL Design

## Collections

There are the following collections in our MongoDB database: discussioncategories, discussionposts, discussiontopics, and users. Compared to the MySQL database, the table subscriptions and userVotes are not represented as separate collections. subscriptions are saved in the user Collection. userVotes are embedded in the discussionposts collection.

### Discussioncategories

The Categories are stored in its own Collection. A Category contains the same data as in the SQL Database. Categories are stored in its own collection because they are needed on their own frequently. Every time a new topic is created all categories are queried. If they were embedded in the discussion topics collection all topics (we assume that there will be much more topics than categories) needed to be queried to find all categories. Our second reason is to make it easier to update or add categories.

Category names should be unique. To ensure this constraint we use an unique index on the name field of a category. As we are sorting by the category name in one of our reports this index will increase the performance of this operation.

```
{
  _id: ObjectId('62b2c6f1abc30fb4b02db9f1'),
  description: 'Local News description',
  name: 'Local News',
}
```

### Users

The Users Collection contains all data, just like the table in the SQL database. Every subscribed category from the User is saved in an array. We decided to save the subscriptions directly in the User collection to prevent infinitely growing arrays. The array can reach a maximum size of the number of categories. If we would have stored it directly in the categories then the array size would be bound by the number of users that would be much greater than the number of categories. The subscribed categories are not directly saved in the document, only their id is saved.

User login data should be mapped to only one user. Therefore an unique index is used to ensure that each user has its unique email address. This index will be useful to speed up the login process as not all documents need to be queried.



```
{
  _id: ObjectId('62b2c6f1abc30fb4b02db9e7'),
  email: 'Levi_Gerhold@yahoo.com',
  password: 'RtrKFGapBGzI7dN',
  subscriptions: [
    ObjectId('62b2c6f1abc30fb4b02db9f1'),
    ObjectId('62b2c6f1abc30fb4b02db9f3'),
    ObjectId('62b2c6f1abc30fb4b02db9f5'),
    ObjectId('62b2c6f1abc30fb4b02db9f8')
  ],
}
```

## Discussiontopics

The Discussiontopics Collection contains all of the SQL database's data. The number of posts is also saved in the topic. Every time a post is added this field will be increased by one for the corresponding topic. We decided against storing references to posts because we anticipate that there will be a large number of posts per topic, which will lead to an infinite growth of the array and, as a result, we may be dangerously close to exceeding the size limit of a document. The count of posts is important for displaying it on the starting page and our first report.

We used several indexes in the Discussiontopics Collection. First we use a decreasing index for the creation time of the article (datetime). This index will be important for our first report to efficiently find all topics that haven been written in after a specified time. An index on the postsCount attribute will help sorting the result. An index on the userId is needed to find all topics written by a user, without needing to check all documents in the collection. This will be important in our second report.

```
{
  _id: ObjectId('62b2c6f1abc30fb4b02dba20'),
  datetime: ISODate('2021-12-23T07:45:59.000Z'),
  discussionCategory: ObjectId('62b2c6f1abc30fb4b02db9fb'),
  text: 'Ad quos sed nemo.\nDucimus odit fugiat.\nProvident quae ut omnis quod.',
  title: 'topic 24',
  userId: '62b2c6f1abc30fb4b02db9e7',
  postsCount: 6,
}
```

```
}
```

## Discussionposts

All data from the discussionPosts table in the SQL database is also included in the Discussionposts Collection. In addition, the discussionThemeId is used to store a reference to the topic. As a result, each post can be assigned to a specific topic. To keep the postsCount in the topic up to date, adding a post for the related topic will increase the postsCount. The userVotes are attached as an array of subdocuments. These contain the same fields as the corresponding table in the SQL database. UserVotes are directly integrated because they are always used in conjunction with the post and thus a separate collection would make little sense in our opinion. Naturally, there may be issues with the maximum document size, but because userVotes don't store lots of information, we don't see this as a major issue.

Because posts are frequently searched for by their discussionThemeId, we use an index to improve query performance. As we query for replies of comments separately the parentPostId also uses an index to improve the query performance for finding all replies of a post.

```
{
```

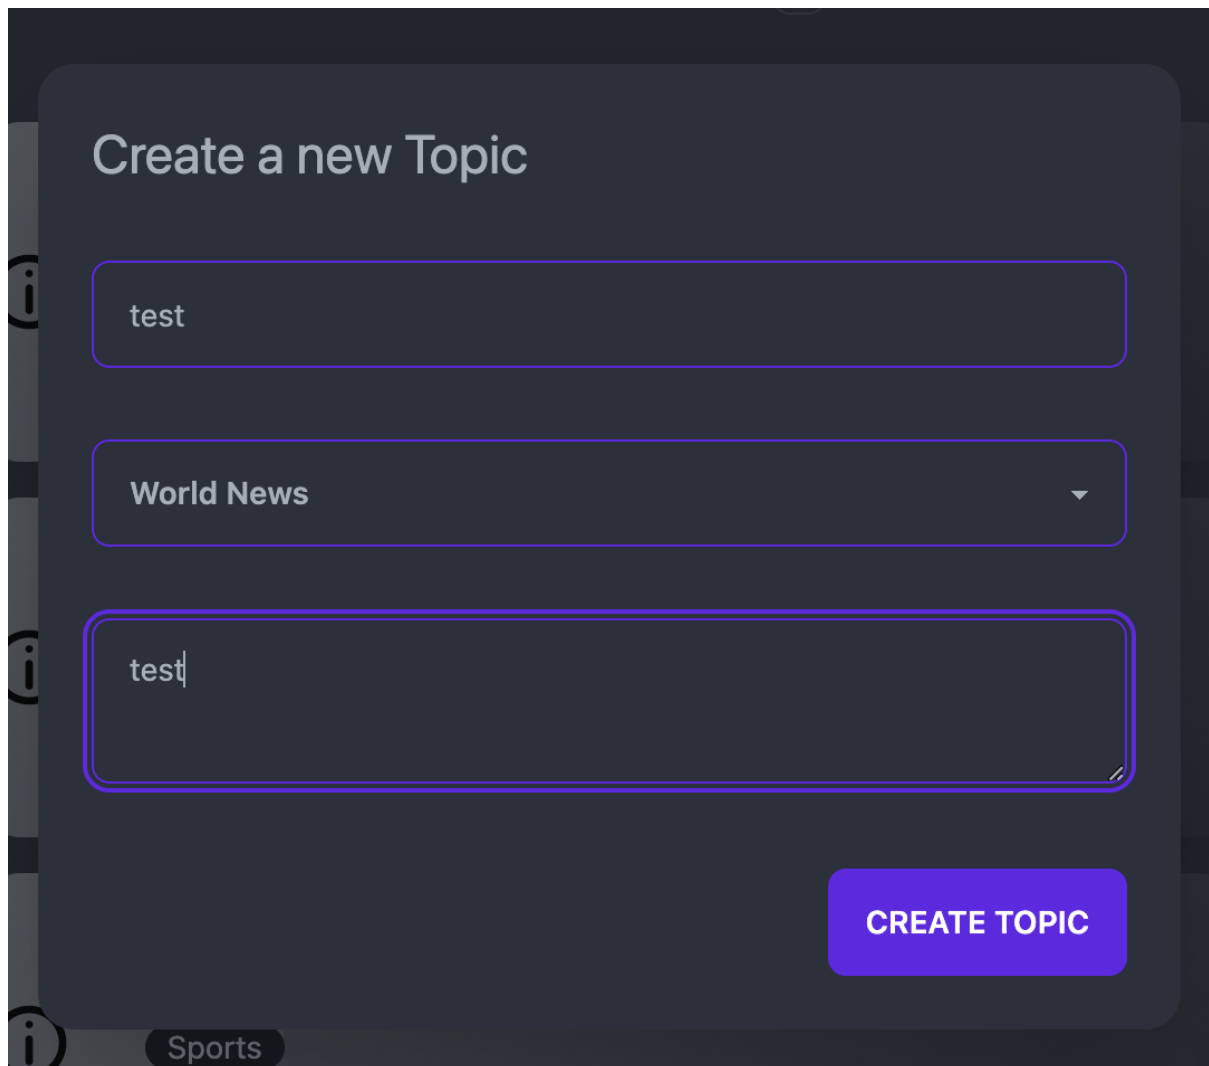
```
  _id: ObjectId('62b2c6f2abc30fb4b02dba48'),
  text: 'Nostrum facere ipsa molestiae molestiae eius veniam quasi.',
  datetime: ISODate('2022-05-11T11:57:21.000Z'),
  discussionThemeId: '62b2c6f1abc30fb4b02dba20',
  parentPostId: null,
  userId: '62b2c6f1abc30fb4b02db9e7',
  userVotes: [
    {
      datetime: ISODate('2022-05-20T00:27:21.000Z'),
      isUpvote: true,
      userId: '62b2c6f1abc30fb4b02db9eb',
      _id: ObjectId('62b2c6f2abc30fb4b02dba49')
    },
    {
      datetime: ISODate('2022-05-12T10:45:40.000Z'),
      isUpvote: true,
      userId: '62b2c6f1abc30fb4b02db9e7',
```

```
      _id: ObjectId('62b2c6f2abc30fb4b02dba4a')
    }
  ],
}
```

## Use-cases and Reports

### Use Case: Creating Discussion Topic

The first main use case is creating topics. This can be done via the topic panel by clicking on “Create topic”. The topic will then be added to the database and displayed on the topic page.



The screenshot shows a dark-themed user interface for creating a new topic. A modal window titled "Create a new Topic" is centered on the screen. It contains three input fields: a text field with the value "test", a dropdown menu currently showing "World News", and a larger text area with the value "test". A blue button labeled "CREATE TOPIC" is positioned at the bottom right of the modal. In the background, a sidebar is partially visible with a "Sports" category selected.

SQL Query:

```
INSERT INTO discussion_topics
(datetime,
id,
text,
title,
userid,
```

```

        discussioncategoryid,

        createdat,

        updatedat)

VALUES      ('2021-06-22 20:25:00',

            DEFAULT,

            'Test text',

            'test',

            2,

            1,

            '2021-06-22 20:25:00',

            '2021-06-22 20:25:00')

```

### MongoDB Query:

```

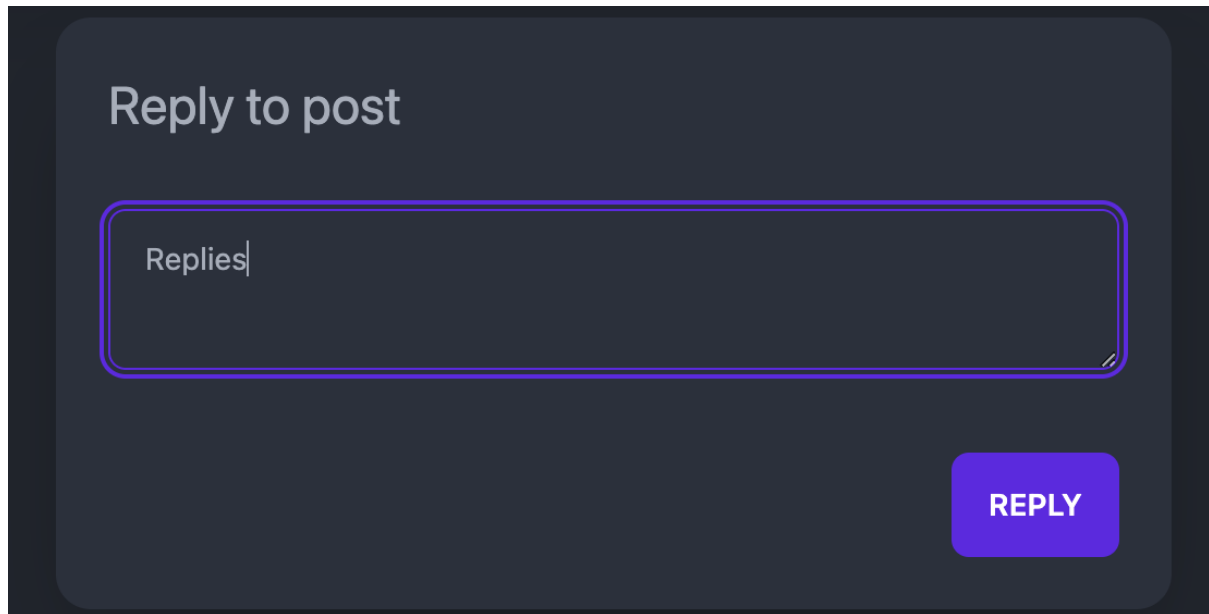
db.discussiontopics.insertOne({
  text: 'some text',
  title: 'title',
  datetime: '2021-06-22 20:25:00',
  userId: '62b2c6f1abc30fb4b02db9e7',
  discussionCategory: '62b2c6f1abc30fb4b02db9fb',
  postsCount: 0
})

```

Both Queries insert the data with a single statement. The same operation is done in the SQL and MongoDB statement. Therefore there will be no performance improvements

## Use Case: Writing a Post

The second main use case is writing a post. This can be done via the post panel by clicking on “Create post” or “reply”. The post will then be added to the database and displayed on the post page.



SQL Query:

```
INSERT INTO discussionposts
    (text,
     datetime,
     id,
     discussionthemeid,
     userid,
     createdat,
     updatedat)
VALUES ('testing',
       '2021-06-22 20:25:00',
       DEFAULT,
       6,
       3,
       '2021-06-22 20:25:00',
       '2021-06-22 20:25:00')
```

## MongoDB Query:

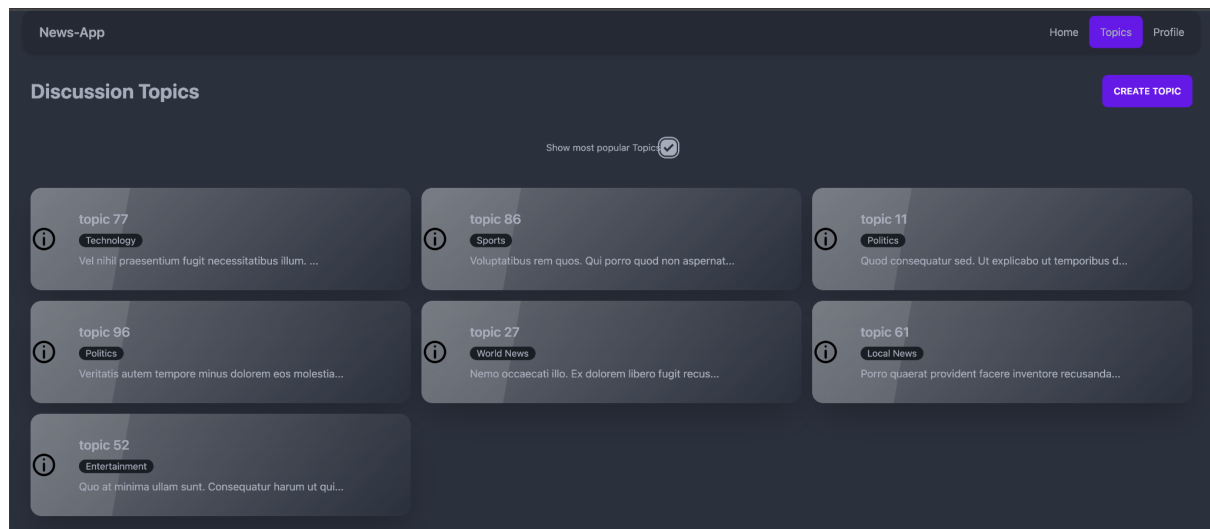
```
db.discussionPosts.insertOne({
  text: 'Nostrum facere ipsa molestiae molestiae eius veniam quasi.',
  datetime: ISODate('2022-05-11T11:57:21.000Z'),
  discussionThemeId: '62b2c6f1abc30fb4b02dba20',
  parentPostId: null,
  userId: '62b2c6f1abc30fb4b02db9e7',
  userVotes: []
})

db.discussiontopics.updateOne({
  _id: '62b2c6f1abc30fb4b02dba20'
}, {
  $inc: {
    postsCount: 1
  }
})
```

Adding Posts in the SQL Database only needs to insert the post in the discussionPosts table, while the mongodb query needs to update the postsCount of the topic and add the post to the collection. Therefore the mongodb insert is more expensive than the SQL insert

## Report: Which are interesting topics of discussions?

This report shows interesting topics that have been written in the last year. Topics are considered as interesting if they get many posts. To show the most interesting first, they are ordered by their number of posts. This report can be called by ticking the checkbox on the topics page.



## SQL Query:

```
SELECT DiscussionTopic.datetime,
       DiscussionTopic.id,
       DiscussionTopic.text,
       DiscussionTopic.title,
       DiscussionTopic.userid,
       DiscussionTopic.discussioncategoryid,
       DiscussionTopic.createdat,
       DiscussionTopic.updatedat,
       Count(discussionposts.id) AS PostCount,
       DiscussionCategory.NAME AS discussioncategory.NAME,
       DiscussionCategory.description AS discussioncategory.description,
       DiscussionCategory.id AS discussioncategory.id,
       DiscussionCategory.createdat AS discussioncategory.createdAt,
       DiscussionCategory.updatedat AS discussioncategory.updatedAt
FROM discussion_topics AS DiscussionTopic
LEFT OUTER JOIN discussioncategories AS DiscussionCategory
    ON DiscussionTopic.discussioncategoryid =
        DiscussionCategory.id
LEFT OUTER JOIN discussionposts AS DiscussionPosts
    ON DiscussionTopic.id = discussionposts.discussionthemeid
WHERE DiscussionTopic.datetime >= '2021-06-22 20:25:00'
GROUP BY DiscussionTopic.id
ORDER BY postcount DESC
```



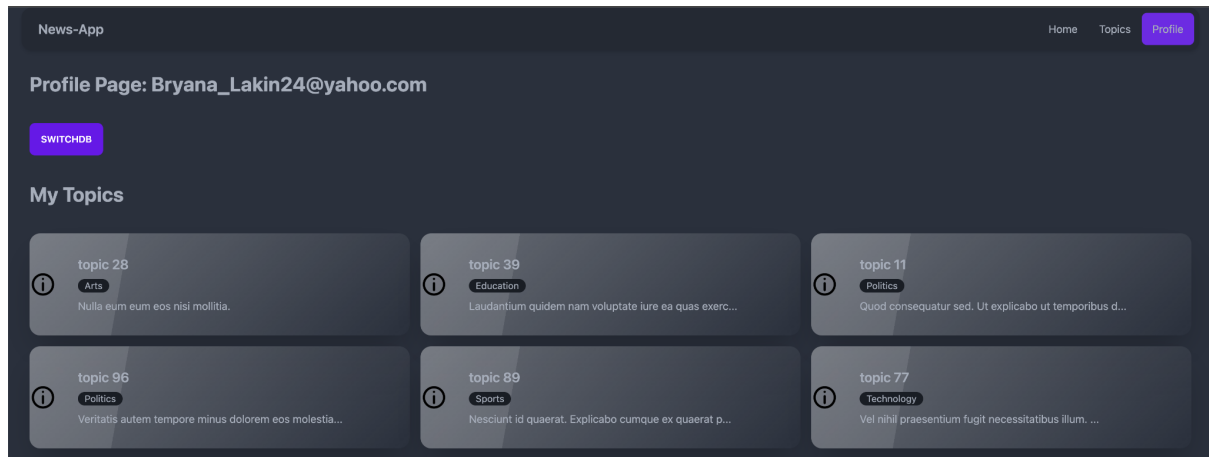
## MongoDB Query:

```
db.discussiontopics.aggregate([{\n  $match: {\n    datetime: {\n      $gte: ISODate('2021-12-23T07:45:59.000Z')\n    }\n  }\n}, {\n  $sort: {\n    postsCount: -1\n  }\n}, {\n  $lookup: {\n    from: 'discussioncategories',\n    localField: 'discussionCategory',\n    foreignField: '_id',\n    as: 'discussionCategory'\n  }\n}, {\n  $unwind: {\n    path: '$discussionCategory'\n  }\n}])
```

The SQL Query needs 2 joins as all data is in different tables. The MongoDB query manages this report with only one lookup operation. The aggregation does not need to lookup (and unwind) the discussionPosts as the number of Posts is stored directly in the topics. Because of our chosen indexes (datetime, postscount) not all documents need to be checked this will lead to a reduction of the query time. The MongoDB query is more performant than the sql query.

## Report: What are all the topics a user has created sorted by category

Our second report shows topics that have been written by a user and sorts them by the name of the category. This report can be called by accessing the profile page.



## SQL Query:

```
SELECT DiscussionTopic.datetime,
       DiscussionTopic.id,
       DiscussionTopic.text,
       DiscussionTopic.title,
       USER.id,
       DiscussionTopic.discussioncategoryid,
       DiscussionTopic.createdat,
       DiscussionTopic.updatedat,
       DiscussionCategory.NAME AS discussioncategory.NAME,
       DiscussionCategory.description AS discussioncategory.description,
       DiscussionCategory.id AS discussioncategory.id,
       DiscussionCategory.createdat AS discussioncategory.createdAt,
       DiscussionCategory.updatedat AS discussioncategory.updatedAt
FROM discussion_topics AS DiscussionTopic
LEFT OUTER JOIN discussioncategories AS DiscussionCategory
    ON DiscussionTopic.discussioncategoryid =
        DiscussionCategory.id
LEFT OUTER JOIN users AS USER
    ON DiscussionTopic.userid = USER.id
WHERE DiscussionTopic.userid = '1'
ORDER BY DiscussionCategory.NAME ASC
```

## MongoDB Query:

```
db.discussiontopics.aggregate([
  $match: {
    userId: '62b38509d6ff45edc0a3614d'
  },
  {
    $lookup: {
      from: 'discussioncategories',
      localField: 'discussionCategory',
      foreignField: '_id',
      as: 'discussionCategory'
    }
  }
])
```

```
}, {
  $unwind: {
    path: '$discussionCategory'
  }
}, {
  $sort: {
    'discussionCategory.name': 1
  }
})
```

For this report the needed joins could not be reduced, but the MongoDB query should be faster by utilizing our indices. As mentioned in the first report by using the index on the `userId` not all topics need to be checked. This will improve the query performance if many topics are stored in the collection.

## **Distribution of work**

Fabian Frauenberger:

- Main use case: Writing a Post
- Report: Which are interesting topics of discussions?
- Data Import
- Data Migration
- SSL Certificates
- Docker

Samuel Mitterrutzner:

- Main use case: Creating Discussion Topic
- Report: What are all the topics a user has created sorted by category
- Frontend
- Docker
- Overall setup (Nx, Next.js and Express)